

```
In [15]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, accuracy_score, precision_score,
```

```
In [28]: df = pd.read_csv('Social_Network_Ads.csv')
```

```
In [29]: df
```

Out[29]:

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0
...
395	15691863	Female	46	41000	1
396	15706071	Male	51	23000	1
397	15654296	Female	50	20000	1
398	15755018	Male	36	33000	0
399	15594041	Female	49	36000	1

400 rows × 5 columns

```
In [30]: df.isna().sum()
df.drop("Gender", axis=1, inplace=True)
```

```
In [31]: # splitting dataset into test and train sets
x = df.iloc[:, :-1].values
y = df.iloc[:, -1].values
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.20, ra
```

In [12]: x

```
Out[12]: array([[15624510, 'Male', 19, 19000],
                [15810944, 'Male', 35, 20000],
                [15668575, 'Female', 26, 43000],
                ...,
                [15654296, 'Female', 50, 20000],
                [15755018, 'Male', 36, 33000],
                [15594041, 'Female', 49, 36000]], dtype=object)
```

```
In [32]: 1 y
```

```
Out[32]: array([0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1,  
                1, 1, 1, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0,  
                0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,  
                0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,  
                0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,  
                0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,  
                0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,  
                0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,  
                0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,  
                0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1,  
                0, 1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 0, 1, 1, 0, 1, 0, 0, 0, 1, 1, 0,  
                1, 1, 0, 1, 0, 1, 0, 1, 0, 0, 1, 1, 0, 1, 0, 0, 1, 1, 0, 1, 1, 0,  
                1, 1, 0, 0, 1, 0, 0, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 0, 1,  
                0, 1, 0, 1, 1, 1, 1, 0, 0, 0, 1, 1, 0, 1, 1, 1, 1, 1, 0, 0, 0, 1,  
                1, 0, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0, 1, 1, 0, 1, 1, 0, 0, 0, 1, 1,  
                0, 1, 0, 0, 1, 0, 1, 0, 0, 1, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0,  
                1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1,  
                0, 1, 0, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1,  
                1, 1, 0, 1], dtype=int64)
```

```
In [33]: # Train logistic regression model on training set
classifier = LogisticRegression(random_state=0)
classifier.fit(x_train, y_train)
```

```
Out[33]: LogisticRegression
LogisticRegression(random_state=0)
```

```
In [40]: y_pred = classifier.predict(x_test)
y_pred_df = pd.DataFrame(y_pred, columns=['Purchase'])
```

In [42]: y_pred_df

Out[42]:

	Purchase
0	0
1	0
2	0
3	0
4	0
...	...
75	0
76	0
77	0
78	0
79	1

80 rows × 1 columns

In [43]: cm = confusion_matrix(y_test, y_pred)
cm

Out[43]: array([[56, 2],
[12, 10]], dtype=int64)

In [44]: *# Extract true positives (TP), false positives (FP), true negatives (TN), false negatives (FN)*
TP = cm[1, 1]
FP = cm[0, 1]
TN = cm[0, 0]
FN = cm[1, 0]

In [45]: accuracy = accuracy_score(y_test, y_pred)
error_rate = 1 - accuracy
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

In [46]: cm

Out[46]: array([[56, 2],
[12, 10]], dtype=int64)

In [47]: TP

Out[47]: 10

In [48]: FP

Out[48]: 2

In [49]: TN

Out[49]: 56

In [50]: FN

Out[50]: 12

In [51]: accuracy

Out[51]: 0.825

In [52]: error_rate

Out[52]: 0.17500000000000004

In [53]: precision

Out[53]: 0.8333333333333334

In [54]: recall

Out[54]: 0.45454545454545453

In [55]: f1

Out[55]: 0.5882352941176471

In []: